

Jiawei Huang

[Email](#)

[GitHub](#)

[Homepage](#)

[Google Scholar](#)

EDUCATION

ETH Zurich

Sept. 2022 – Now

Ph.D. Student in Computer Science

University of Illinois at Urbana-Champaign

Jan. 2021 – Aug. 2022

Ph.D. Student in Computer Science

Awarded by *Master's Degree in Computer Science* on Aug. 2022

(*Ph.D. study was interrupted because of visa issue ([P.P. 10043](#)).)*

Beihang University, Beijing, China

Sept. 2015 – July 2019

Undergraduate Student in Computer Science

Awarded by *Bachelor's Degree in Computer Science* on July 2019

INTERNSHIP EXPERIENCES

Reinforcement Learning | University of Illinois at Urbana-Champaign | Visiting Scholar | July 2019 – July 2020

- **Doubly Robust Policy Gradient Method**
- **Offline Reinforcement Learning**

Reinforcement Learning + Robotics | Tecent Robotics-X Lab | Research Intern | Aug. 2020 – Jan. 2021

- **Robot Motion Imitation with Reinforcement Learning and Imitation Learning**

Reinforcement Learning | Microsoft Research Asia | Research Intern | Summer 2021

- **Deployment-Efficient Reinforcement Learning**
- **Reinforcement Learning with Tiered Structure**

PUBLICATION

1. [Preprint] **Robust Knowledge Transfer in Tiered Reinforcement Learning** ([arXiv](#))
Jiawei Huang, Niao He
2. [NeurIPS 2022] **Tiered Reinforcement Learning: Pessimism in the Face of Uncertainty and Constant Regret** ([arXiv](#))
Jiawei Huang, Li Zhao, Tao Qin, Wei Chen, Nan Jiang, Tie-Yan Liu
3. [ICML 2022 **Long Oral**] **A Minimax Learning Approach to Off-Policy Evaluation in Confounded POMDP** ([arXiv](#))
Chengchun Shi, Masatoshi Uehara, **Jiawei Huang**, Nan Jian
4. [ICLR 2022 **Spotlight**] **Towards Deployment-Efficient Reinforcement Learning: Lower Bound and Optimality** ([OpenReview](#))
Jiawei Huang, Jinglin Chen, Li Zhao, Tao Qin, Nan Jiang, Tie-Yan Liu
5. [AISTATS 2022] **On the Convergence Rate of Density-Ratio Based Off-Policy Policy Gradient** ([arXiv](#))
Jiawei Huang, Nan Jiang
6. [NeurIPS 2020] **Minimax Confidence Interval for Off-Policy Evaluation and Policy Optimization** ([arXiv](#))
Nan Jiang, **Jiawei Huang**

7. [ICML 2020] From Importance Sampling to Doubly Robust Policy Gradient ([arXiv](#))
Jiawei Huang, Nan Jiang
8. [ICML 2020] Minimax Weight and Q-Function Learning for Off-Policy Evaluation ([arXiv](#))
Masatoshi Uehara, Jiawei Huang, Nan Jiang
9. [ECCV 2020] WeightNet: Revisiting the Design Space of Weight Networks ([arXiv](#))
Ningning Ma, Xiangyu Zhang, Jiawei Huang, Jian Sun

PROGRAMMING SKILLS.

Python (TensorFlow, PyTorch), Java, C/C++.